

# Exploring the Depths of Sound:

## An Analysis of CNNs for Environmental Sound Classification

HUMAN DATA ANALYTICS PROJECT PRESENTATION



ELAHEH AHMADIESLAMLOO

GEORGE P. PRODAN

# INTRODUCTION

## Environmental Sound Classification



### Motivation

With the rise of audio-based technologies the creation of efficient algorithms for ESC is becoming a crucial issue

### Objective

Our goal is to implement CNN models to solve the ESC problem

### Contributions

- Performance study of different architectures
- Model scalings
- Study of regularization techniques


# Why CNN?



**Various new data data augmentation techniques**

**Extract features**

**Successful performance in a variety of audio classification tasks**



# RELATED WORK

## Convolutional Neural Networks


Study	Method	Dataset	Accuracy
Lee	Multiscale CNN	DCASE2017	63.4%
Piczak	Temporal CNN + attention	ESC-10	72.2%
Abdoli	1D CNN	UrbanSound8k	89%
Zhang	CNN-Gated Recurrent Unit	ESC-10	92.3%
Zhang	CNN-Gated Recurrent Unit	ESC-50	87.43%
Zhang	CNN-Gated Recurrent Unit	UrbanSound8K	96.1%

## Residual Layers ← ResNet Architectures

S. R. Kaiming He, Xiangyu Zhang and J. Sun, *Deep residual learning for image recognition*, 2015.

## Recurrent CNNs

J. Sang, S. Park, and J. Lee, *Convolutional recurrent neural networks for urban sound classification using raw waveforms*, 2018.



# Data augmentation

**Improve the performance**


**Learn more patterns (not in original)**

**Techniques:**

- ❖ **Speed changing**
- ❖ **pitch shifting**
- ❖ **time shifting**
- ❖ **Clipping**
- ❖ **time masking, frequency masking, inserting Gaussian noise**

# Main steps

1. **Data Pre-processing(mono,44100 hz)**
2. **Data Loading(subclass of tf.keras.Sequence )**
3. **Model Training (40-100 epochs )**
4. **Model Testing**



# AUDIO REPRESENTATIONS


## Raw audio waveforms

the raw waveforms are loaded using *librosa* library

mono audio & 44.1 kHz

## Feature extraction :

- Mel spectrogram
- Mel-frequency Cepstral Coefficients (MFCCs)



# DATASET: ESC-50

## ESC-50

50 classes of sounds from various environments

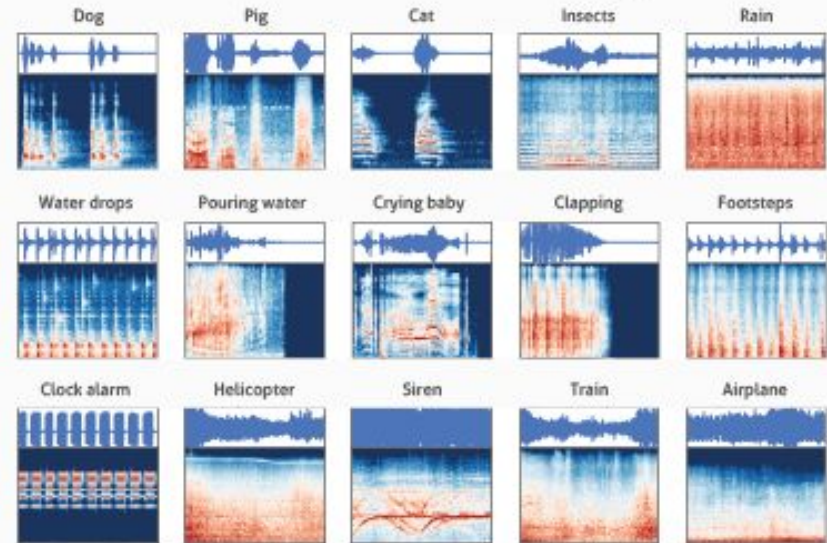
5 second clips, 40 clips/class ➡ 2000 clips

Split the dataset into 80/10/10

### Stratified split

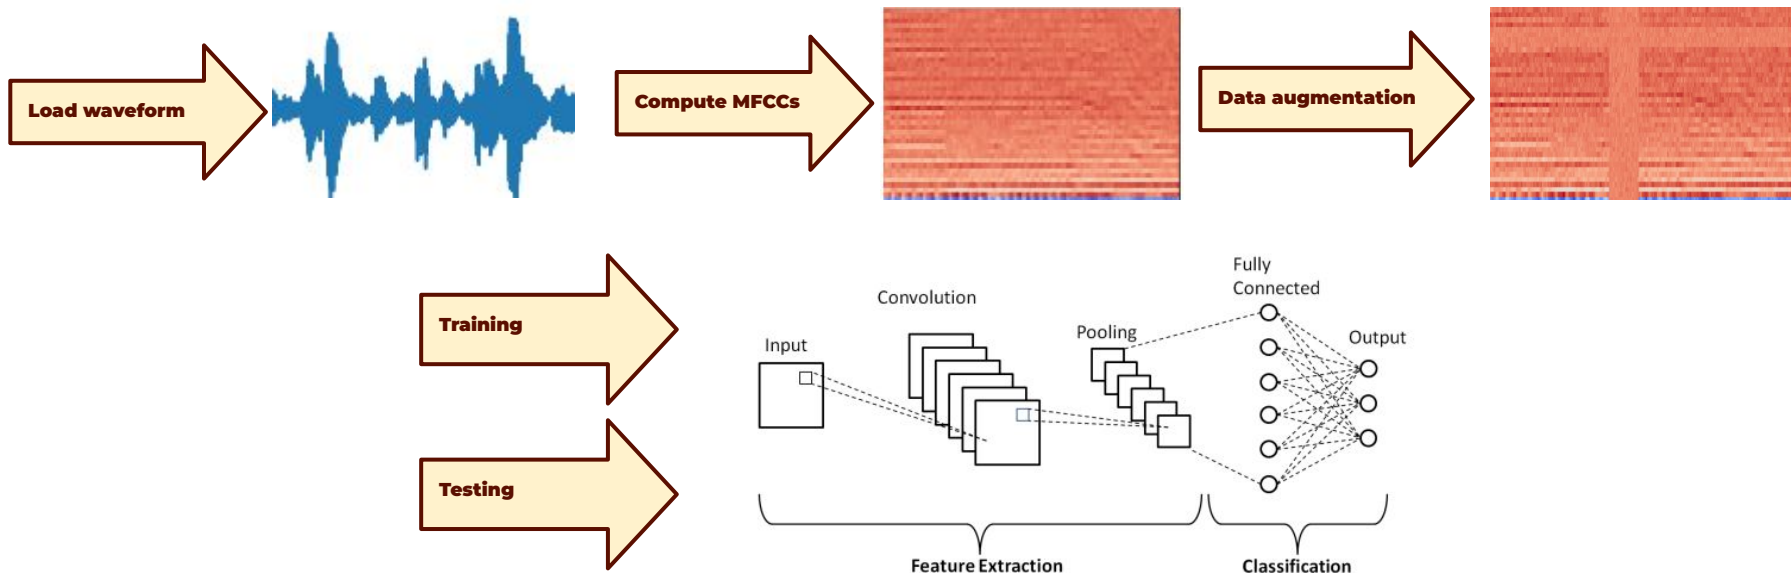
class distribution in the three sets is representative of the class distribution in the whole dataset

Overview of selected classes of recordings

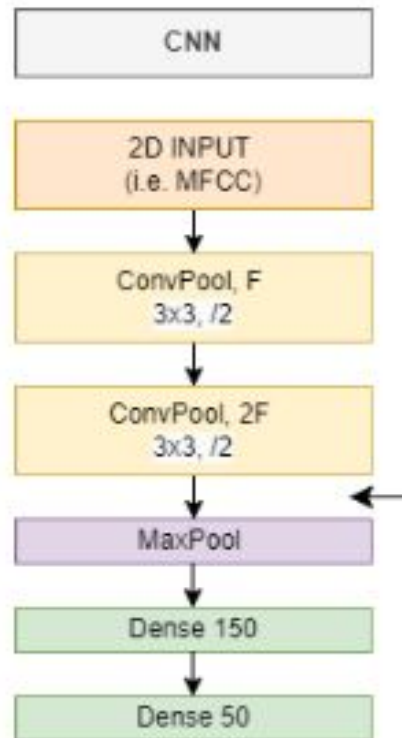




# LEARNING FRAMEWORK



# Basic CNN



**extra layers added here**

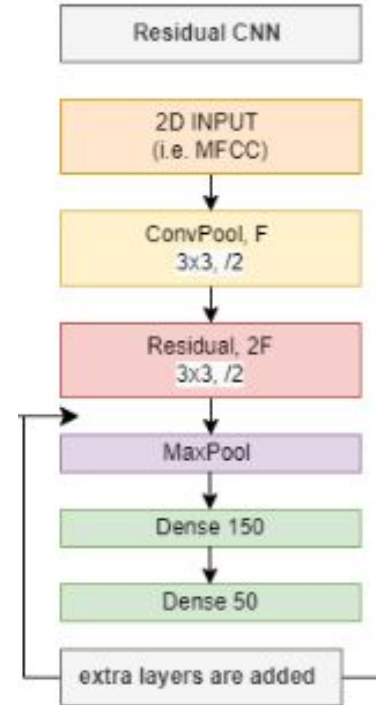
**F = number of filters**

# CNN + Residual Layers

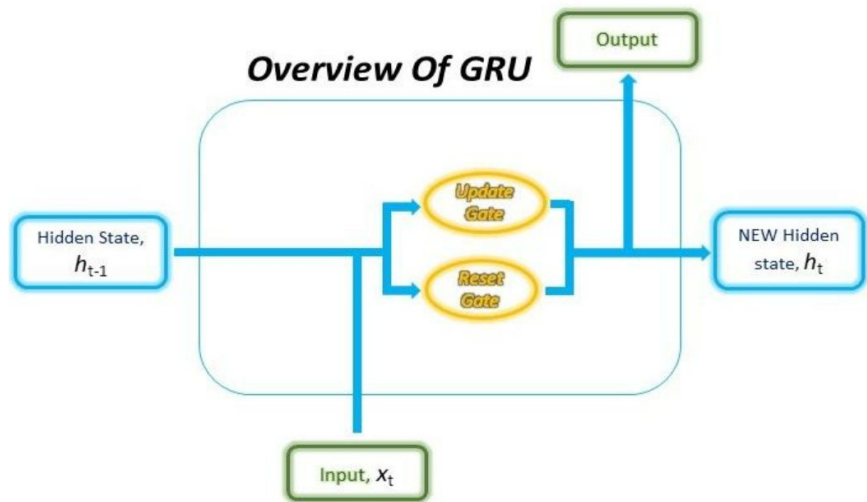
## Residual Layers

significant advancement in deep neural networks for image processing tasks

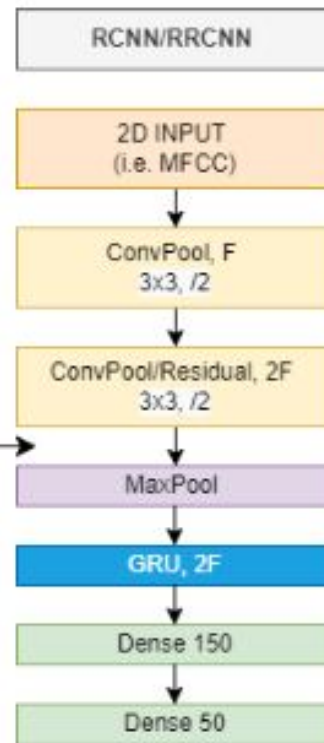
solved the problem of "gradient vanishing" in deep networks with many stacked layers



# Recurrent CNNs



**extra layers added here**



# FIRST EXPERIMENTS

performed on a CNN with 2 layers

## Kernel Size

3x3	5x5	7x7
88.5%	84.5%	87.5%

## Pool Size

2x2	4x4	8x8
88.5%	84%	79%

## Width Scaling

filters	test acc, %	param. (M)	best epoch / 50
4	86.00	0.5	50
8	86.50	1.0	35
16	90.50	2.0	25
32	88.50	4.0	29
64	89.50	8.0	35
128	90.00	16.2	23
256	87.00	33.1	18

# Trying Residual Layers, or GRUs

## Test Accuracy (2L)

<b>CNN</b>	<b>88.5%</b>
<b>CNN+Residual Layers</b>	<b>86.5%</b>
<b>CNN+GRU</b>	<b>93.0%</b>
<b>CNN+Residual Layers+GRU</b>	<b>91.5%</b>

# DEPTH SCALING

<b>Architecture</b>	<b>2L</b>	<b>3L</b>	<b>4L</b>	<b>5L</b>	<b>6L</b>
<b>CNN</b>	88.5%	90.0%	91.5%	90.5%	91.0%
<b>CNN+Residual Layers</b>	86.5%	91.5%	87.5%	89.0%	89.0%
<b>CNN+GRU</b>	93.0%	92.0%	88.5%	92.0%	89.0%
<b>CNN+Residual Layers+GRU</b>	91.5%	92.0%	93.5%	92.0%	92.0%

# REGULARIZATION TECHNIQUES


- **Dropout Layers**
- **Data Augmentation**
  - Time Masking
  - Frequency Masking
  - Gaussian Noise

## Grid search results


- Using both frequency and time masks
- 4% gaussian noise
- 0.5 dropout



## without regularization



## with regularization



## BEST RESULTS

Network	Input	Test Acc., %
Our best CNN	MFCC	91.5
Our best ResCNN	MFCC	91.5
Our best RCNN	MFCC	93.0
Our best RRCNN	MFCC	93.5
SoundNet [21]	Raw	87.5
Pre-trained transformer* [22]	Mel	95.7



# WORST RESULTS

(on a CNN with 2 layers)

**Not using BN: 79.5%**

**Using larger LR:**

**2% for 0.1**

**15% for 0.01**

**Using only one FC layer: 10.5%**

# DEMO



```
class CNN(tf.keras.Model):
    def __init__(self, F=32, n=2, num_classes=50, p_dropout=None, residual=False):
        super().__init__()
        self.F = F
        self.n = n
        self.p_dropout = p_dropout
        self.residual = residual

        Layer = ResBlock2D if residual else ConvPool2D

        activation = 'relu'
        self.description = f'CNN_{n}layers'
        hidden_units = 150

        if residual:
            self.identity_1 = Identity(F)
            if n > 2:
                self.identity_2 = Identity(F)

        # 1st block
        self.layer_1 = ConvPool2D(F)

        self.layer_2 = Layer(F*2, maxpool=True)
        if n > 2:
            if residual: hidden_units = 500
            self.layer_3 = Layer(F*2, maxpool=True)
        if n > 3:
            hidden_units = 450
            self.layer_4 = Layer(F*2, maxpool=False)

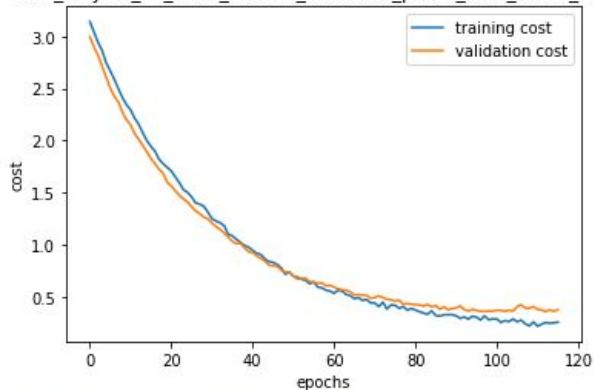
        if n > 4:
            self.layer_5 = Layer(F*2, maxpool=False, padding="same")
        if n > 5:
            self.layer_6 = Layer(F*2, maxpool=False, padding="same")
        k = 10 if n == 2 else 5
        self.pool_gru = tf.keras.layers.MaxPool2D((k, 1), strides=(k, 1), padding='same')
```

# DEMO



```
[5.770000 min, epoch 114] Train Cost: 0.202372 --- Validation Cost: 0.371020 --- Train accuracy: 0.931073 --- Validation accuracy: 0.893000
[5.830000 min, epoch 115] Train Cost: 0.222256 --- Validation Cost: 0.382603 --- Train accuracy: 0.963750 --- Validation accuracy: 0.915000
[5.880000 min, epoch 116] Train Cost: 0.240868 --- Validation Cost: 0.378268 --- Train accuracy: 0.960625 --- Validation accuracy: 0.910000
[5.920000 min, epoch 117] Train Cost: 0.256871 --- Validation Cost: 0.363395 --- Train accuracy: 0.946875 --- Validation accuracy: 0.915000
[5.980000 min, epoch 118] Train Cost: 0.250522 --- Validation Cost: 0.378364 --- Train accuracy: 0.951250 --- Validation accuracy: 0.920000
[6.040000 min, epoch 119] Train Cost: 0.253783 --- Validation Cost: 0.364822 --- Train accuracy: 0.948125 --- Validation accuracy: 0.920000
[6.090000 min, epoch 120] Train Cost: 0.260722 --- Validation Cost: 0.380681 --- Train accuracy: 0.953125 --- Validation accuracy: 0.910000
```

CNN\_2layers\_50\_E120\_OAdam\_LR0.0001\_pD0.5\_A0.4\_feb13\_t2240



[TESTING] Test accuracy: 0.930000 --- Test cost: 0.268969

# DEMO



TIME	MODEL NAME	INPUT	BATCH	EPOCHS	AUGMENT PRO	DROPOUT PRO	OPTIMIZER	LR	TRAIN ACC	VAL ACC	TEST ACC	BEST EPOCH	PARAMETERS	FILTERS	RESIDUAL	TIME/EPOCH
02-13 13:45	CNN_6layersGR	mfcc		50	100	0	Adam	0.0001	1	0.93	0.92	87	6267086	128	TRUE	4.3742
02-13 13:57	CNN_6layersGR	mfcc		50	100	0	Adam	0.0001	1	0.915	0.89	69	3210702	128	FALSE	3.0054
02-13 14:12	CNN_5layersGR	mfcc		50	100	0	Adam	0.0001	1	0.925	0.92	61	2619598	128	FALSE	2.8041
02-13 14:28	CNN_4layersGR	mfcc		50	100	0	Adam	0.0001	1	0.905	0.885	60	2028494	128	FALSE	2.8686
02-13 14:45	CNN_3layersGR	mfcc		50	100	0	Adam	0.0001	1	0.92	0.92	79	1329940	128	FALSE	2.882
02-13 15:15	CNN_3layersGR	mfcc		50	100	0	Adam	0.0001	1	0.92	0.915	82	1329940	128	FALSE	2.754
02-13 15:21	CNN_2layersGR	mfcc		50	100	0	Adam	0.0001	1	0.93	0.93	100	738836	128	FALSE	2.4178
02-13 15:33	CNN_2layers	mfcc		50	100	0	Adam	0.0001	0.9988	0.915	0.865	53	3004436	128	TRUE	2.2961
02-13 15:44	CNN_2layers	mfcc		50	60	0	Adam	0.0001	0.9994	0.92	0.885	60	2302484	128	FALSE	1.347833333
02-13 16:20	CNN_3layers	mfcc		50	60	0	Adam	0.0001	1	0.92	0.915	30	5653710	128	TRUE	1.723166667
02-13 16:25	CNN_4layers	mfcc		50	60	0	Adam	0.0001	0.9956	0.93	0.875	28	6835918	128	TRUE	1.878833333
02-13 16:30	CNN_5layers	mfcc		50	60	0	Adam	0.0001	0.9994	0.94	0.89	25	8018126	128	TRUE	2.098666667
02-13 16:43	CNN_6layers	mfcc		50	60	0	Adam	0.0001	0.9956	0.915	0.885	26	9200334	128	TRUE	2.355333333

# CONCLUSIONS

- **Features such as BNs and GRUs improve the performance**
- **It is not always helpful to have more layers**
- **Data Augmentation can improve the results by 1-2%**
- **Residual Layers without GRU are not a better choice**

**Thank you for your attention!**